

```
In [3]: import numpy as np
import pandas as pd

import seaborn as sns
import plotly.express as px
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings('ignore')

%matplotlib inline
```

```
In [4]: df=pd.read_csv("C:/Users/sindh/Downloads/archive (6)/WineQT.csv")
```

```
In [5]: df.head()
```

Out[5]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alco
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	

```
In [6]: print(df.columns)
print(df.shape)
```

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual suga
r',
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'densit
y',
      'pH', 'sulphates', 'alcohol', 'quality', 'Id'],
      dtype='object')
(1143, 13)
```

In [7]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1143 entries, 0 to 1142
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1143 non-null   float64
1   volatile acidity       1143 non-null   float64
2   citric acid            1143 non-null   float64
3   residual sugar         1143 non-null   float64
4   chlorides              1143 non-null   float64
5   free sulfur dioxide    1143 non-null   float64
6   total sulfur dioxide   1143 non-null   float64
7   density                1143 non-null   float64
8   pH                    1143 non-null   float64
9   sulphates              1143 non-null   float64
10  alcohol                1143 non-null   float64
11  quality                1143 non-null   int64
12  Id                     1143 non-null   int64
dtypes: float64(11), int64(2)
memory usage: 116.2 KB

```

In [8]: df.describe().T

Out[8]:

	count	mean	std	min	25%	50%	75%	
fixed acidity	1143.0	8.311111	1.747595	4.60000	7.10000	7.90000	9.100000	15.9
volatile acidity	1143.0	0.531339	0.179633	0.12000	0.39250	0.52000	0.640000	1.5
citric acid	1143.0	0.268364	0.196686	0.00000	0.09000	0.25000	0.420000	1.0
residual sugar	1143.0	2.532152	1.355917	0.90000	1.90000	2.20000	2.600000	15.5
chlorides	1143.0	0.086933	0.047267	0.01200	0.07000	0.07900	0.090000	0.6
free sulfur dioxide	1143.0	15.615486	10.250486	1.00000	7.00000	13.00000	21.000000	68.0
total sulfur dioxide	1143.0	45.914698	32.782130	6.00000	21.00000	37.00000	61.000000	289.0
density	1143.0	0.996730	0.001925	0.99007	0.99557	0.99668	0.997845	1.0
pH	1143.0	3.311015	0.156664	2.74000	3.20500	3.31000	3.400000	4.0
sulphates	1143.0	0.657708	0.170399	0.33000	0.55000	0.62000	0.730000	2.0
alcohol	1143.0	10.442111	1.082196	8.40000	9.50000	10.20000	11.100000	14.9
quality	1143.0	5.657043	0.805824	3.00000	5.00000	6.00000	6.000000	8.0
Id	1143.0	804.969379	463.997116	0.00000	411.00000	794.00000	1209.500000	1597.0

```
In [9]: df.nunique()
```

```
Out[9]: fixed acidity      91  
volatile acidity    135  
citric acid         77  
residual sugar      80  
chlorides           131  
free sulfur dioxide  53  
total sulfur dioxide 138  
density             388  
pH                  87  
sulphates           89  
alcohol             61  
quality             6  
Id                  1143  
dtype: int64
```


```
In [10]: df.duplicated().sum()
```

```
Out[10]: 0
```

```
In [11]: sns.set(style="whitegrid")  
print(df['quality'].value_counts())  
fig = plt.figure(figsize = (10,6))  
sns.countplot(data=df, x='quality', palette='pastel')
```

```
quality  
5    483  
6    462  
7    143  
4     33  
8     16  
3      6  
Name: count, dtype: int64
```

```
Out[11]: <Axes: xlabel='quality', ylabel='count'>
```

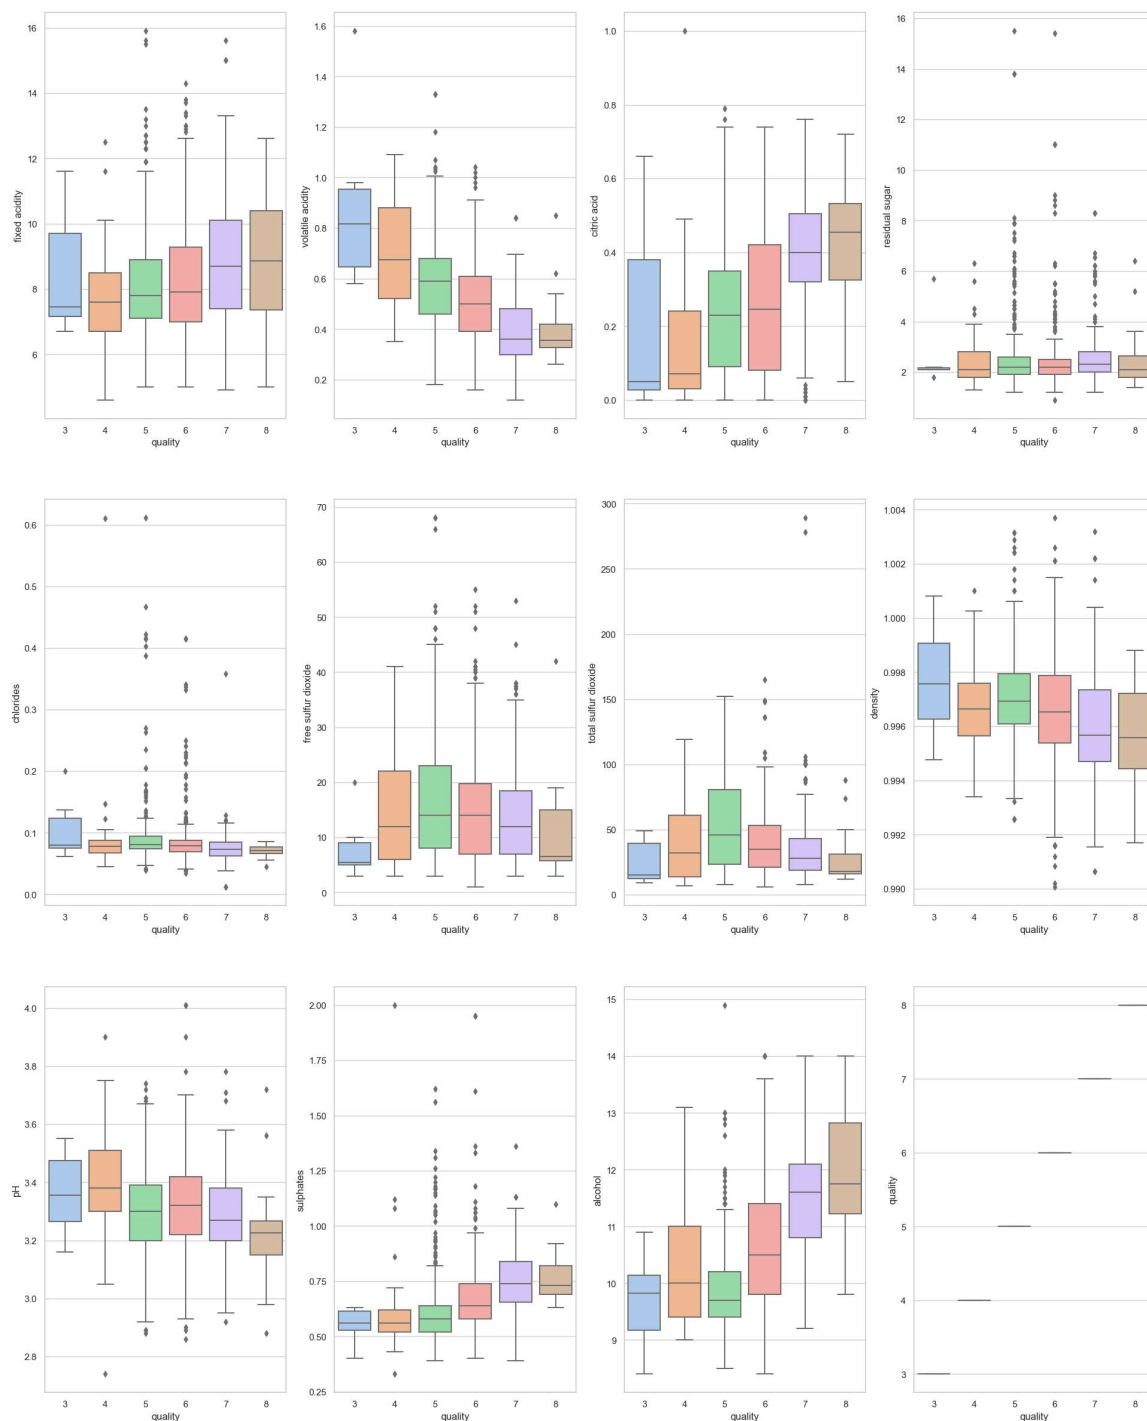


```

In [12]: import warnings
warnings.filterwarnings("ignore")


sns.set(style="whitegrid")
fig, ax1 = plt.subplots(3,4, figsize=(24,30))
k = 0
columns = list(df.columns)
for i in range(3):
    for j in range(4):
        sns.boxplot(data=df, x='quality', y=columns[k], ax=ax1[i][j], palette=
            k += 1
plt.show()

```



```
In [13]: plt.figure(figsize = (15,15))
sns.heatmap(df.corr(),annot=True, cmap= 'PuBuGn')
```

Out[13]: <Axes: >

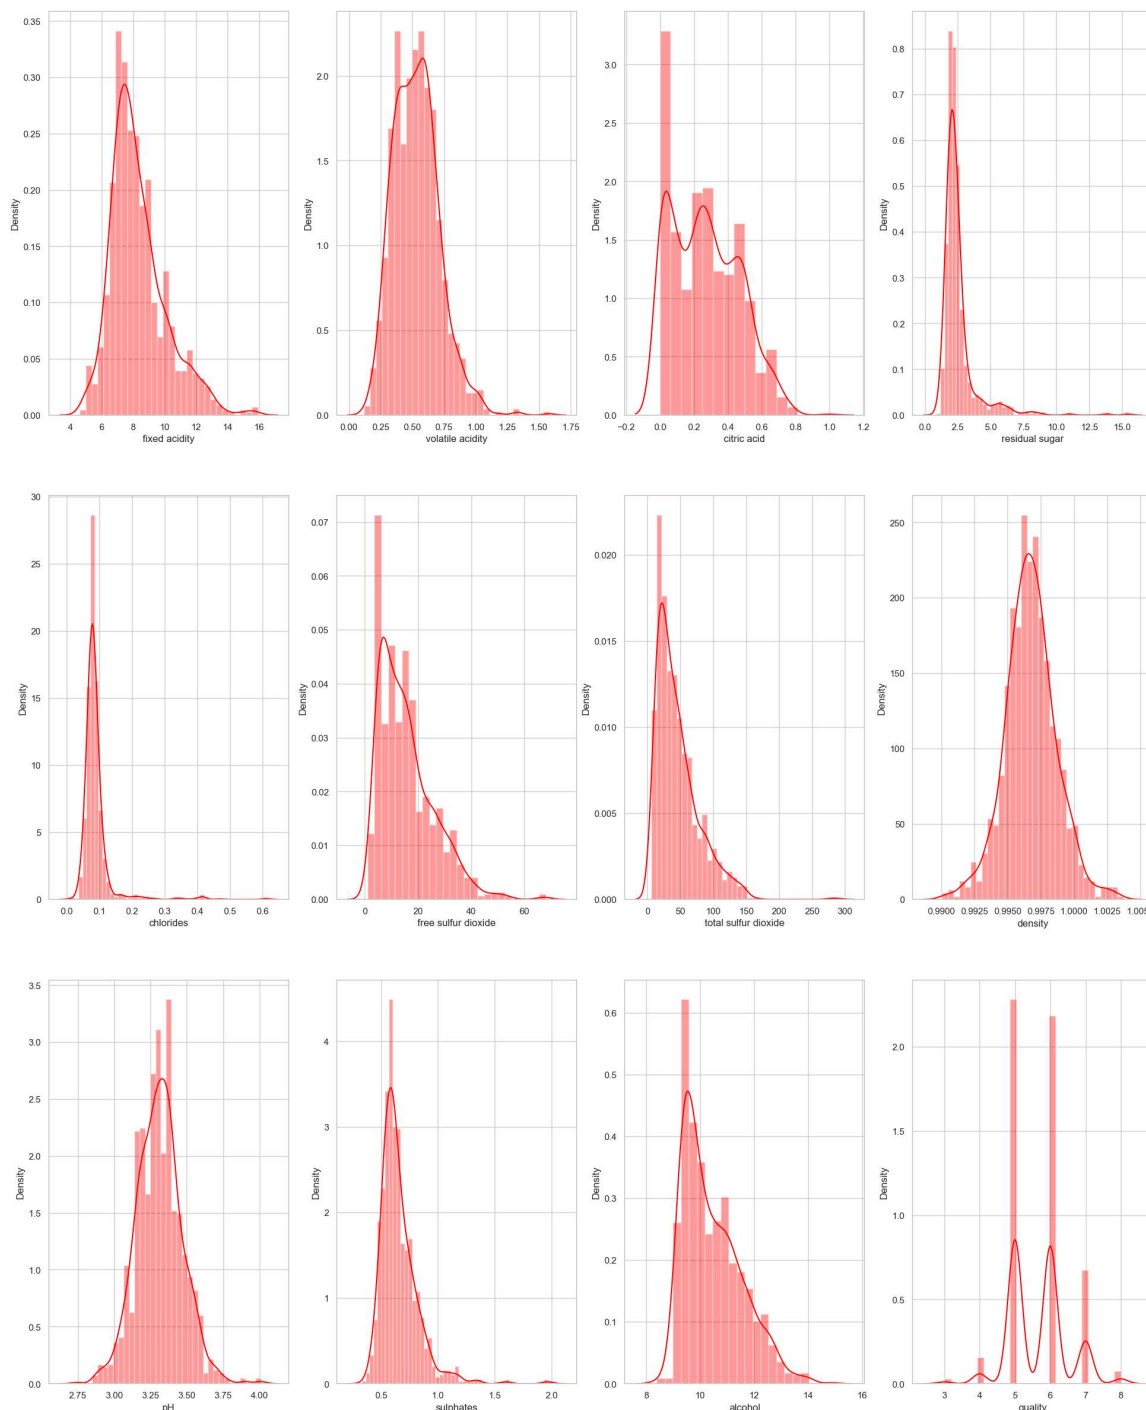


```

In [14]: color = sns.color_palette("pastel")

fig, ax1 = plt.subplots(3,4, figsize=(24,30))
k = 0
columns = list(df.columns)
for i in range(3):
    for j in range(4):
        sns.distplot(df[columns[k]], ax = ax1[i][j], color = 'red')
        k += 1
plt.show()

```

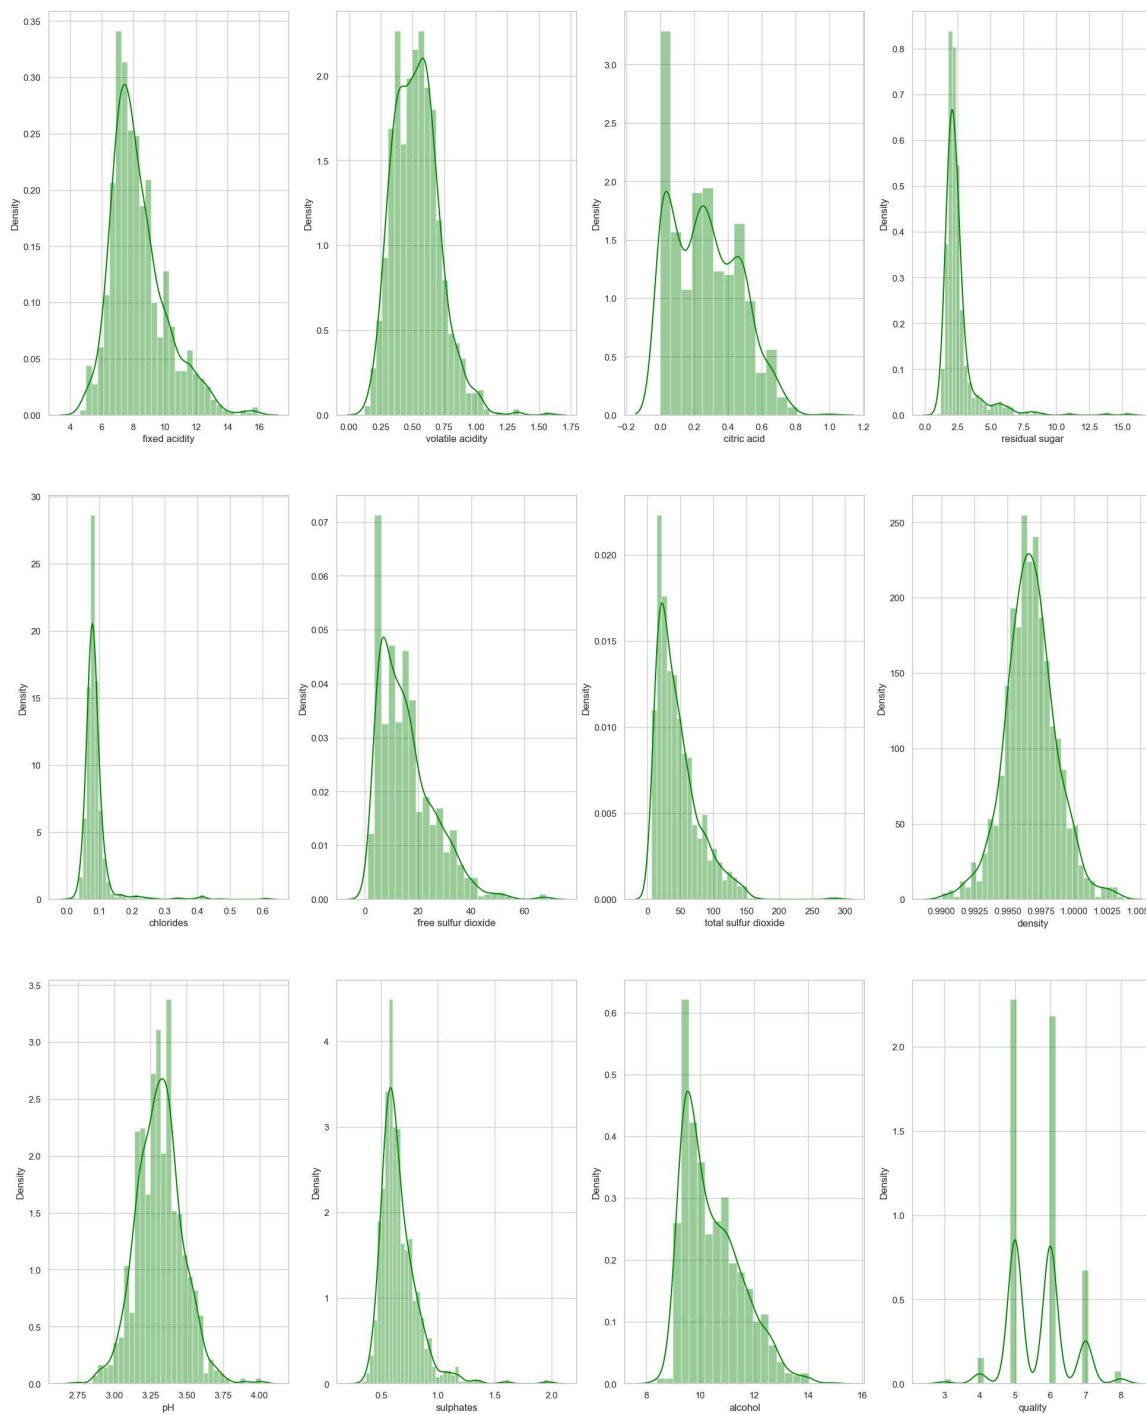


```

In [15]: color = sns.color_palette("pastel")

fig, ax1 = plt.subplots(3,4, figsize=(24,30))
k = 0
columns = list(df.columns)
for i in range(3):
    for j in range(4):
        sns.distplot(df[columns[k]], ax = ax1[i][j], color = 'green')
        k += 1
plt.show()

```



```
In [16]: def log_transform(col):
          return np.log(col[0])

df['residual sugar'] = df[['residual sugar']].apply(log_transform, axis=1)
df['chlorides'] = df[['chlorides']].apply(log_transform, axis=1)
df['free sulfur dioxide'] = df[['free sulfur dioxide']].apply(log_transform, axis=1)
df['total sulfur dioxide'] = df[['total sulfur dioxide']].apply(log_transform, axis=1)
df['sulphates'] = df[['sulphates']].apply(log_transform, axis=1)
```

```
In [17]: df.corr()['quality'].sort_values(ascending=False)
```

```
Out[17]: quality          1.000000
         alcohol         0.484866
         sulphates       0.315097
         citric acid      0.240821
         fixed acidity     0.121970
         Id              0.069708
         residual sugar    0.031487
         pH              -0.052453
         free sulfur dioxide -0.054185
         total sulfur dioxide -0.170128
         density          -0.175208
         chlorides        -0.175391
         volatile acidity  -0.407394
         Name: quality, dtype: float64
```

```
In [18]: df_3 = df[df.quality==3]
         df_4 = df[df.quality==4]
         df_5 = df[df.quality==5]
         df_6 = df[df.quality==6]
         df_7 = df[df.quality==7]
         df_8 = df[df.quality==8]
```

```
In [19]: from sklearn.utils import resample

df_3_upsampled = resample(df_3, replace=True, n_samples=600, random_state=123)
df_4_upsampled = resample(df_4, replace=True, n_samples=600, random_state=123)
df_7_upsampled = resample(df_7, replace=True, n_samples=600, random_state=123)
df_8_upsampled = resample(df_8, replace=True, n_samples=600, random_state=123)
df_5_downsampled = df[df.quality==5].sample(n=600, replace=True).reset_index()
df_6_downsampled = df[df.quality==6].sample(n=600, replace=True).reset_index()
```

```
In [20]: Balanced_df = pd.concat([df_3_upsampled, df_4_upsampled, df_7_upsampled,
                                   df_8_upsampled, df_5_downsampled, df_6_downsampled])


Balanced_df.quality.value_counts()
```

```
Out[20]: quality
         3      600
         4      600
         7      600
         8      600
         5      600
         6      600
         Name: count, dtype: int64
```




```
In [24]: plt.figure(figsize=(10,6))  
sns.countplot(x='quality', data=Balanced_df, order=[3, 4, 5, 6, 7, 8], palet
```

```
Out[24]: <Axes: xlabel='quality', ylabel='count'>
```




```
In [25]: plt.figure(figsize = (12,6))  
sns.barplot(x='quality', y = 'alcohol', data = df, palette = 'coolwarm')
```

```
Out[25]: <Axes: xlabel='quality', ylabel='alcohol'>
```



```
In [26]: plt.figure(figsize=(15,15))
Balanced_df.corr().quality.apply(lambda x: abs(x)).sort_values(ascending=False)
plt.title("Top 10 highly correlated features", size=20, pad=26)
plt.xlabel("Correlation coefficient")
plt.ylabel("Features")
```

```
Out[26]: Text(0, 0.5, 'Features')
```



```
In [27]: selected_features = ['fixed acidity', 'volatile acidity', 'citric acid', 'chlorides',
                             'free sulfur dioxide', 'total sulfur dioxide', 'density',
                             'sulphates', 'alcohol']
X = Balanced_df[selected_features]
y = Balanced_df.quality
```

```
In [28]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X, y, test_size=0.3,random_
```

```
In [29]: from sklearn.neighbors import KNeighborsClassifier
for n_neighbors in [5,10,15,20]:
    model = KNeighborsClassifier(n_neighbors)
    model.fit(X_train, y_train)
    scr = model.score(X_test, y_test)
    print("For n_neighbors = ", n_neighbors, ", score is ",scr)
```

```
For n_neighbors = 5  score is  0.8277777777777777
For n_neighbors = 10  score is  0.7638888888888888
For n_neighbors = 15  score is  0.7240740740740741
For n_neighbors = 20  score is  0.6657407407407407
```

```
In [30]: KNN_Model = KNeighborsClassifier(n_neighbors=5, weights='distance')
KNN_Model.fit(X_train, y_train)
results = KNN_Model.fit(X_train, y_train)
```

```
In [31]: KNN_train_predictions = KNN_Model.predict(X_train)
```

```
In [32]: KNN_test_predictions = KNN_Model.predict(X_test)
```

```
In [33]: from sklearn.metrics import classification_report, confusion_matrix

print("\n Train Data: KNN_Confusion Matrix:\n ")
print(confusion_matrix(y_train, KNN_train_predictions))

print("\n Train Data: KNN_Classification Report:\n ")
print(classification_report(y_train, KNN_train_predictions))

print("\n \n Test Data: KNN_Confusion Matrix: \n ")
print(confusion_matrix(y_test, KNN_test_predictions))

print("\n Test Data: KNN_Classification Report:\n ")
print(classification_report(y_test, KNN_test_predictions))
```

Train Data: KNN_Confusion Matrix:

```
[[422  0  0  0  0  0]
 [  0 392  0  0  0  0]
 [  0  0 423  0  0  0]
 [  0  0  0 436  0  0]
 [  0  0  0  0 423  0]
 [  0  0  0  0  0 424]]
```

Train Data: KNN_Classification Report:

	precision	recall	f1-score	support
3	1.00	1.00	1.00	422
4	1.00	1.00	1.00	392
5	1.00	1.00	1.00	423
6	1.00	1.00	1.00	436
7	1.00	1.00	1.00	423
8	1.00	1.00	1.00	424
accuracy			1.00	2520
macro avg	1.00	1.00	1.00	2520
weighted avg	1.00	1.00	1.00	2520

Test Data: KNN_Confusion Matrix:

```
[[178  0  0  0  0  0]
 [  0 208  0  0  0  0]
 [  1  3 147 19  7  0]
 [  1  3 11 133 15  1]
 [  0  0  0  6 171  0]
 [  0  0  0  0  0 176]]
```

Test Data: KNN_Classification Report:

	precision	recall	f1-score	support
3	0.99	1.00	0.99	178
4	0.97	1.00	0.99	208
5	0.93	0.83	0.88	177
6	0.84	0.81	0.83	164
7	0.89	0.97	0.92	177
8	0.99	1.00	1.00	176
accuracy			0.94	1080
macro avg	0.94	0.93	0.93	1080
weighted avg	0.94	0.94	0.94	1080

In []:

